

April 15, 2024

```
[1]: import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[3]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
[5]: data = pd.DataFrame(boston.data)
data.head()
data.columns = boston.feature_names
data.head()
```

```
[5]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```
[6]: data['MEDV'] = boston.target
```

```
[8]: data.shape
```

```
[8]: (506, 14)
```

```
[9]: data.columns
```

```
[9]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
          'PTRATIO', 'B', 'LSTAT', 'MEDV'],
          dtype='object')
```

```
[13]: data.dtypes
```

```
[13]: CRIM      float64
      ZN      float64
      INDUS   float64
      CHAS     float64
      NOX     float64
      RM      float64
      AGE     float64
      DIS     float64
      RAD     float64
      TAX     float64
      PTRATIO  float64
      B       float64
      LSTAT   float64
      MEDV    float64
      dtype: object
```

```
[14]: data.isnull().sum
```

```
[14]: <bound method DataFrame.sum of          CRIM      ZN  INDUS   CHAS    NOX      RM
AGE    DIS    RAD    TAX  \
0      False  False  False  False  False  False  False  False  False  False  False
1      False  False  False  False  False  False  False  False  False  False  False
2      False  False  False  False  False  False  False  False  False  False  False
3      False  False  False  False  False  False  False  False  False  False  False
4      False  False  False  False  False  False  False  False  False  False  False
..      ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
501     False  False  False  False  False  False  False  False  False  False  False
502     False  False  False  False  False  False  False  False  False  False  False
503     False  False  False  False  False  False  False  False  False  False  False
504     False  False  False  False  False  False  False  False  False  False  False
505     False  False  False  False  False  False  False  False  False  False  False

      PTRATIO      B  LSTAT  MEDV
0      False  False  False  False
1      False  False  False  False
2      False  False  False  False
3      False  False  False  False
4      False  False  False  False
..      ...    ...    ...    ...
501     False  False  False  False
502     False  False  False  False
```

```

503    False  False  False  False
504    False  False  False  False
505    False  False  False  False

```

```
[506 rows x 14 columns]>
```

```
[15]: data[data.isnull().any(axis = 1)]
```

```

[15]: Empty DataFrame
      Columns: [CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT, MEDV]
      Index: []

```

```
[16]: data.describe()
```

```

[16]:
      count    CRIM      ZN      INDUS      CHAS      NOX      RM  \
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean    3.613524   11.363636   11.136779    0.069170    0.554695    6.284634
std     8.601545   23.322453    6.860353    0.253994    0.115878    0.702617
min     0.006320    0.000000    0.460000    0.000000    0.385000    3.561000
25%     0.082045    0.000000    5.190000    0.000000    0.449000    5.885500
50%     0.256510    0.000000    9.690000    0.000000    0.538000    6.208500
75%     3.677083   12.500000   18.100000    0.000000    0.624000    6.623500
max     88.976200  100.000000   27.740000    1.000000    0.871000    8.780000

      count    AGE      DIS      RAD      TAX      PTRATIO      B  \
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean    68.574901    3.795043    9.549407   408.237154   18.455534   356.674032
std    28.148861    2.105710    8.707259  168.537116    2.164946   91.294864
min     2.900000    1.129600    1.000000  187.000000   12.600000    0.320000
25%    45.025000    2.100175    4.000000  279.000000   17.400000   375.377500
50%    77.500000    3.207450    5.000000  330.000000   19.050000   391.440000
75%    94.075000    5.188425   24.000000  666.000000   20.200000   396.225000
max   100.000000   12.126500   24.000000  711.000000   22.000000   396.900000

      count    LSTAT      MEDV
count  506.000000  506.000000
mean    12.653063   22.532806
std     7.141062    9.197104
min     1.730000    5.000000
25%     6.950000   17.025000
50%    11.360000   21.200000
75%    16.955000   25.000000
max    37.970000   50.000000

```

```

[19]: corr = data.corr()
      corr.shape

```

```
plt.figure(figsize=(20, 20))
sns.heatmap(corr, cbar = True, square = True, fmt = '.1f', annot = True,
            annot_kws = {'size':15})
```

[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1d5bb5927c0>



```
[21]: x = data.drop(['MEDV'], axis = 1)
      y = data['MEDV']
```

```
[22]: from sklearn.model_selection import train_test_split
```

```
[23]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.3,
↳ random_state = 4)
```

```
[24]: from sklearn.linear_model import LinearRegression
```

```
[25]: lm = LinearRegression()
```

```
[26]: lm.fit(X_train, y_train)
```

```
[26]: LinearRegression()
```

```
[28]: lm.intercept_
```

```
[28]: 36.357041376595205
```

```
[35]: coefficients= pd.DataFrame([X_train.columns, lm.coef_]).T
coefficients = coefficients.rename(columns = {0: 'Attribute', 1:
↳ 'Coefficients'})
coefficients
```

```
[35]:   Attribute Coefficients
0      CRIM      -0.12257
1        ZN      0.0556777
2     INDUS -0.00883428
3      CHAS      4.69345
4      NOX     -14.4358
5       RM      3.28008
6      AGE  -0.00344778
7      DIS     -1.55214
8      RAD      0.32625
9      TAX    -0.0140666
10  PTRATIO    -0.803275
11       B      0.00935369
12    LSTAT    -0.523478
```

```
[37]: y_pred = lm.predict(X_train)
```


```
[39]: print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_train,
↳ y_pred))*(len(y_train)-1)/
(len(y_train)-X_train.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
```


```
Adjusted R^2: 0.736910342429894
```

MAE: 3.08986109497113
MSE: 19.07368870346903
RMSE: 4.367343437774162

```
[40]: plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```




```
[41]: plt.scatter(y_pred, y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



```
[42]: sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```

C:\Users\Tej\anaconda3\lib\site-packages\seaborn\distributions.py:2551:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).

```
warnings.warn(msg, FutureWarning)
```



```
[43]: y_test_pred = lm.predict(X_test)

acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_test, y_test_pred)) * (len(y_test) - 1) / (len(y_test) - X_test.shape[1] - 1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7121818377409195
Adjusted R^2: 0.6850685326005713
MAE: 3.8590055923707407
MSE: 30.053993307124127
RMSE: 5.482152251362974
```